

Importing required libraries

```
In [1]:
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

Stroke Dataset: Loading

- Stroke prediction dataset (gathered from kaggle. Author: Fedesoriano. Name: Stroke-prediction-dataset)
- The purpose of the dataset is to forecast the possibility of an individual having a stroke, utilising input parameters such as gender, age, presence of different ailments, and smoking status. Each row of data provides pertinent information about the patient.
- The dataset includes 5110 rows and 12 columns, with each row representing a different patient and each column representing a different attribute or feature. The columns in the dataset are:

Variable	Description
iD	Unique identifier for each patient
gender	Gender of the patient (Male/Female/Other)
age	Age of the patient in years
hypertension	Whether the patient has hypertension or not (0 - No, 1 - Yes)
heart_disease	Whether the patient has heart disease or not (0 - No, 1 - Yes)
ever_married	Whether the patient has ever been married or not (Yes/No)
work_type	Type of occupation of the patient (Private, Self-employed, Govt_job, children, Never_worked)
Residence_type	Area type of residence of the patient (Urban/Rural)
avg_glucose_level	Average glucose level in the patient's blood
bmi	Body mass index of the patient
smoking_status	Habitual infromation, patient's smoking status (formerly smoked, never smoked, smokes, unknown)
stroke	Whether the patient has had a stroke or not (0 - No, 1 - Yes)

```
In [2]:
stroke_data = pd.read_csv("healthcare-dataset-stroke-data.csv")
```

```
In [3]:
stroke_data.head(10)
```

Out[3]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	strok
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked	
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	NaN	never smoked	
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoked	
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smokes	
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never smoked	
5	56669	Male	81.0	0	0	Yes	Private	Urban	186.21	29.0	formerly smoked	
6	53882	Male	74.0	1	1	Yes	Private	Rural	70.09	27.4	never smoked	
7	10434	Female	69.0	0	0	No	Private	Urban	94.39	22.8	never smoked	
8	27419	Female	59.0	0	0	Yes	Private	Rural	76.15	NaN	Unknown	
9	60491	Female	78.0	0	0	Yes	Private	Urban	58.57	24.2	Unknown	

Checking the summary of statistics for each column of the DataFrame

In [4]:

```
stroke_data.describe()
```

Out[4]:

	id	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000	5110.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893237	0.048728
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.854067	0.215320
min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000	0.000000
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.500000	0.000000
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.100000	0.000000
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	33.100000	0.000000
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000	1.000000

In [5]:

```
stroke_data.info() # this is to obtain basic information about a DataFrame, and is often used as a first step in expl
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    5110 non-null   int64
 1   gender                5110 non-null   object
 2   age                  5110 non-null   float64
 3   hypertension          5110 non-null   int64
 4   heart_disease         5110 non-null   int64
 5   ever_married          5110 non-null   object
 6   work_type             5110 non-null   object
 7   Residence_type        5110 non-null   object
 8   avg_glucose_level     5110 non-null   float64
 9   bmi                   4909 non-null   float64
10   smoking_status        5110 non-null   object
11   stroke                5110 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
```

In [6]:

```
# Calculating the mean of the 'bmi' column
mean_value = stroke_data ['bmi'].mean()
stroke_data ['bmi'].fillna(value = mean_value, inplace = True)
```

In [7]:

```
stroke_data.head(10)
# it can be seen that all the missing values have been replaced by the mean
```

Out[7]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.600000	formerly smoked
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.893237	never smoked
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.500000	never smoked
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.400000	smokes
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.000000	never smoked
5	56669	Male	81.0	0	0	Yes	Private	Urban	186.21	29.000000	formerly smoked
6	53882	Male	74.0	1	1	Yes	Private	Rural	70.09	27.400000	never smoked
7	10434	Female	69.0	0	0	No	Private	Urban	94.39	22.800000	never smoked
8	27419	Female	59.0	0	0	Yes	Private	Rural	76.15	28.893237	Unknown
9	60491	Female	78.0	0	0	Yes	Private	Urban	58.57	24.200000	Unknown

Exploratory Data Analysis (EDA)

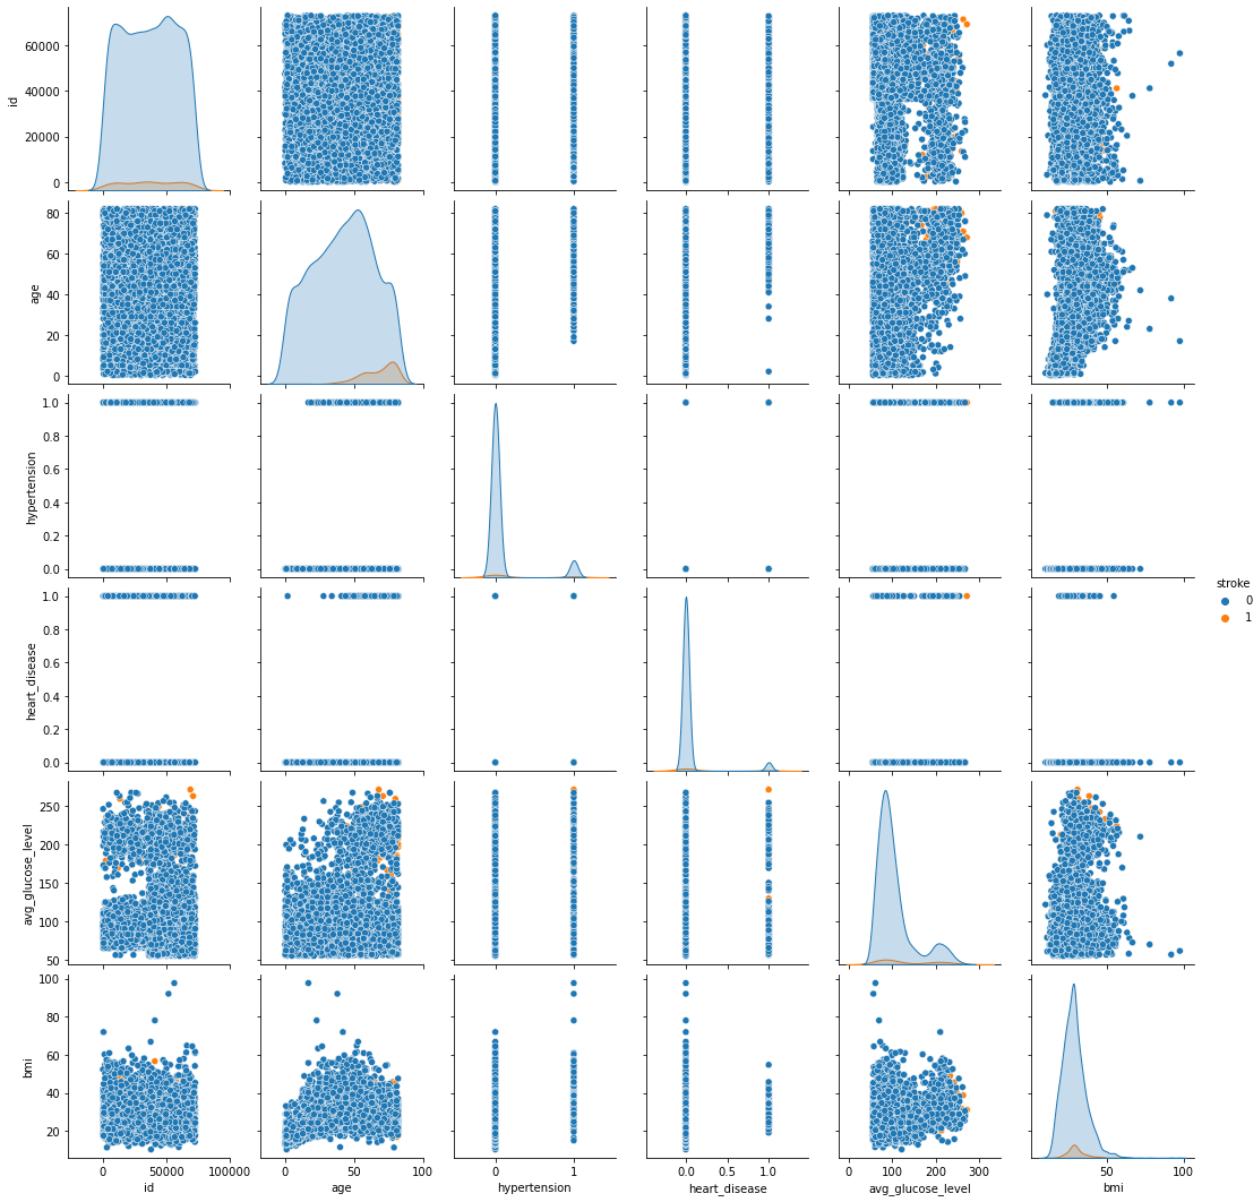
Calculating the mean inside the 'bmi' attribute and replacing all the null and missing values with the mean

In [8]:

sns.pairplot(stroke\_data, hue="stroke")

Out[8]:

<seaborn.axisgrid.PairGrid at 0x7f7ba3b51040>



Distribution by Gender


In [9]:

```
# Count the occurrences of each gender
gender_count = stroke_data['gender'].value_counts()

# Plotting a Bar chart
plt.bar(gender_count.index, gender_count.values)

# Adding the labels and the char title
plt.xlabel('Gender')
plt.ylabel('Count')
plt.title('Gender Count')

plt.show()
```



In [10]:

```
gender_count
```

Out[10]:


```
Female    2994
Male      2115
Other         1
Name: gender, dtype: int64
```

### Distribution by Smoking Status

In [11]:

```
#Pie chart of stroke occurrence by smoking status
smoking_stroke = stroke_data[stroke_data["stroke"] == 1]["smoking_status"].value_counts()
plt.pie(smoking_stroke, labels=smoking_stroke.index, autopct='%1.1f%%', startangle=90)
plt.title("Pie Chart of Stroke Occurrence by Smoking Status")
plt.show()
```

Pie Chart of Stroke Occurrence by Smoking Status



### Distribution by Heart Disease

In [12]:

```
# Count the occurrences of each heart disease status
hd_count = stroke_data['heart_disease'].value_counts()

# Plotting a bar chart
plt.bar(['No Heart Disease', 'With Heart Disease'], hd_count.values, color=['#9bb7d4', '#0f4c81'])

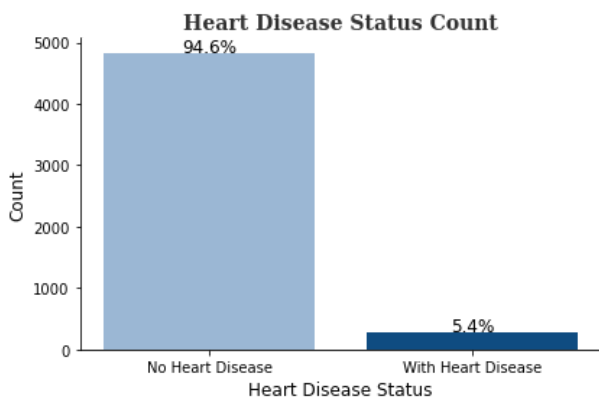
# Adding the labels and the chart title
plt.xlabel('Heart Disease Status', fontsize=12)
plt.ylabel('Count', fontsize=12)
plt.title('Heart Disease Status Count', fontsize=14, fontweight='bold', fontfamily='serif', color="#323232")

# Adding the percentage values above the bars
for i, count in enumerate(hd_count.values):
    plt.text(i, count+10, str(round((count/len(stroke_data))*100, 1)) + '%', ha='center', fontsize=12)

# Removing the top and right spines
plt.gca().spines['top'].set_visible(False)
plt.gca().spines['right'].set_visible(False)

# Adjusting the layout
plt.tight_layout()

plt.show()
```




### Distribution of Frequency by Age

In [13]:

```
#this is to check stroke frequency by age

stroke_positive = stroke_data[stroke_data['stroke'] == 1]
stroke_negative = stroke_data[stroke_data['stroke'] == 0]

# Plot the age distributions for stroke and non-stroke patients
plt.hist(stroke_negative['age'], bins=30, alpha=0.6, label='Non-Stroke Frequency')
plt.hist(stroke_positive['age'], bins=20, alpha=0.5, label='Stroke Frequency')
plt.title('Age Distribution for Stroke and Non-Stroke Patients')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.legend(loc='upper left')
plt.show()
```




### Distribution by Stroke and Glucose Levels

In [14]:

```
#Box plot of glucose levels by stroke occurrence
sns.boxplot(x="stroke", y="avg_glucose_level", data=stroke_data)
```


Out[14]:

```
<AxesSubplot:xlabel='stroke', ylabel='avg_glucose_level'>
```



In [15]:

```
#Scatter plot of glucose level and BMI by age:
plt.scatter(stroke_data["avg_glucose_level"], stroke_data["bmi"], c=stroke_data["age"], cmap="viridis")
plt.xlabel("Glucose Level")
plt.ylabel("BMI")
plt.title("Scatter Plot of Glucose Level and BMI by Age")
plt.colorbar()
plt.show()
```




### Distribution Gender and BMI scores

In [16]:

```
#Violin plot of BMI by stroke occurrence and gender
sns.violinplot(x="gender", y="bmi", hue="stroke", data=stroke_data, split=True)
```

Out[16]:


```
<AxesSubplot:xlabel='gender', ylabel='bmi'>
```



### Distribution frequency of Stroke occurrences by Age

In [17]:

```
#Line plot of stroke occurrence by age
sns.lineplot(x="age", y="stroke", data=stroke_data)
plt.xlabel("Age")
plt.ylabel("Stroke Occurrence")
plt.title("Line Plot of Stroke Occurrence by Age")
plt.show()
```




### Distribution of Stroke patients filtered by Hypertension

In [18]:

```
# Count the number of patients with and without hypertension who have had a stroke
hypertension_stroke = stroke_data.loc[(stroke_data['hypertension'] == 1) & (stroke_data['stroke'] == 1)].shape[0]
no_hypertension_stroke = stroke_data.loc[(stroke_data['hypertension'] == 0) & (stroke_data['stroke'] == 1)].shape[0]

# Create a bar chart to visualize the results
sns.set_style('whitegrid')
plt.figure(figsize=(8,6))
ax = sns.barplot(x=['Hypertension', 'No Hypertension'], y=[hypertension_stroke, no_hypertension_stroke], palette=['#F08080', '#FFD700'])
plt.title('Number of Patients with and without Hypertension who have had a Stroke', fontsize=14, fontweight='bold')
plt.xlabel('Hypertension Status', fontsize=12, fontweight='bold')
plt.ylabel('Number of Patients', fontsize=12, fontweight='bold')
plt.xticks(fontsize=12, fontweight='bold')
plt.yticks(fontsize=12, fontweight='bold')
for p in ax.patches:
    ax.annotate(format(p.get_height(), '.0f'), (p.get_x() + p.get_width() / 2., p.get_height()), ha = 'center', va = 'bottom')
plt.show()
```



### Dataset Pre-Processing

#### Dropping unwanted variables

In [19]:

```
# Removing the 'id' column as it is redundant in this case
stroke_data = stroke_data.drop('id', axis=1)
```

In [20]:

```
#This is to take of the decimal value in age
stroke_data['age'] = stroke_data['age'].astype(int)
```

In [21]:

```
stroke_data.head()
```

Out[21]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	Male	67	0	1	Yes	Private	Urban	228.69	36.600000	formerly smoked	1
1	Female	61	0	0	Yes	Self-employed	Rural	202.21	28.893237	never smoked	1
2	Male	80	0	1	Yes	Private	Rural	105.92	32.500000	never smoked	1
3	Female	49	0	0	Yes	Private	Urban	171.23	34.400000	smokes	1
4	Female	79	1	0	Yes	Self-employed	Rural	174.12	24.000000	never smoked	1

### Replacing the non-numerical categorical values with the numerical integers

In [22]:

```
# Creating a dictionary to map non-numerical values with the numerical integers
gender_mapping = {'Male': 0, 'Female':1, 'Other':2} # this is to replace the gender values from 'Male', 'Female' to '
# replace non-numerical values with numerical values using the mapping dictionary
stroke_data['gender'] = stroke_data['gender'].replace(gender_mapping)
stroke_data.head()
```

Out[22]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	0	67	0	1	Yes	Private	Urban	228.69	36.600000	formerly smoked	1
1	1	61	0	0	Yes	Self-employed	Rural	202.21	28.893237	never smoked	1
2	0	80	0	1	Yes	Private	Rural	105.92	32.500000	never smoked	1
3	1	49	0	0	Yes	Private	Urban	171.23	34.400000	smokes	1
4	1	79	1	0	Yes	Self-employed	Rural	174.12	24.000000	never smoked	1

In [23]:

```
# Replcing the 'ever_married' values from non numerical to numeriactal integers
marr_mapping = {'Yes':0, 'No':1}
stroke_data['ever_married'] = stroke_data['ever_married'].replace(marr_mapping)
```

In [24]:

```
stroke_data.head()
```

Out[24]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	0	67	0	1	0	Private	Urban	228.69	36.600000	formerly smoked	1
1	1	61	0	0	0	Self-employed	Rural	202.21	28.893237	never smoked	1
2	0	80	0	1	0	Private	Rural	105.92	32.500000	never smoked	1
3	1	49	0	0	0	Private	Urban	171.23	34.400000	smokes	1
4	1	79	1	0	0	Self-employed	Rural	174.12	24.000000	never smoked	1



In [25]:

```
# Replcing the 'work_type' values from non numerical to numeriactal integers
stroke_data['work_type'].unique()
```

Out[25]:

```
array(['Private', 'Self-employed', 'Govt_job', 'children', 'Never_worked'],
      dtype=object)
```

In [26]:

```
work_mapping = {'Private':0, 'Self-employed':1, 'Govt_job':2, 'children':3, 'Never_worked':4}
stroke_data['work_type'] = stroke_data['work_type'].replace(work_mapping)
stroke_data.head()
```

Out[26]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	0	67	0	1	0	0	Urban	228.69	36.600000	formerly smoked	1
1	1	61	0	0	0	1	Rural	202.21	28.893237	never smoked	1
2	0	80	0	1	0	0	Rural	105.92	32.500000	never smoked	1
3	1	49	0	0	0	0	Urban	171.23	34.400000	smokes	1
4	1	79	1	0	0	1	Rural	174.12	24.000000	never smoked	1

In [27]:

```
# Replcing the 'Residence_type' values from non numerical to numeriactal integers
stroke_data['Residence_type'].unique()
```

Out[27]:

```
array(['Urban', 'Rural'], dtype=object)
```

In [28]:

```
res_mapping = {'Urban':0, 'Rural':1}
stroke_data['Residence_type'] = stroke_data['Residence_type'].replace(res_mapping)
stroke_data.head()
```

Out[28]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	0	67	0	1	0	0	0	228.69	36.600000	formerly smoked	1
1	1	61	0	0	0	1	1	202.21	28.893237	never smoked	1
2	0	80	0	1	0	0	1	105.92	32.500000	never smoked	1
3	1	49	0	0	0	0	0	171.23	34.400000	smokes	1
4	1	79	1	0	0	1	1	174.12	24.000000	never smoked	1

In [29]:

```
# Replcing the 'smoking_status' values from non numerical to numeriactal integers
stroke_data['smoking_status'].unique()
```

Out[29]:

```
array(['formerly smoked', 'never smoked', 'smokes', 'Unknown'],
      dtype=object)
```

In [30]:

```
res_mapping = {'formerly smoked':0, 'never smoked':1, 'smokes':2, 'Unknown':3}
stroke_data['smoking_status'] = stroke_data['smoking_status'].replace(res_mapping)
```

In [31]:

```
# Now all the non-numerical categorical values are replaced with Integers
stroke_data.head(10)
```

Out[31]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	0	67	0	1	0	0	0	228.69	36.600000	0	1
1	1	61	0	0	0	1	1	202.21	28.893237	1	1
2	0	80	0	1	0	0	1	105.92	32.500000	1	1
3	1	49	0	0	0	0	0	171.23	34.400000	2	1
4	1	79	1	0	0	1	1	174.12	24.000000	1	1
5	0	81	0	0	0	0	0	186.21	29.000000	0	1
6	0	74	1	1	0	0	1	70.09	27.400000	1	1
7	1	69	0	0	1	0	0	94.39	22.800000	1	1
8	1	59	0	0	0	0	1	76.15	28.893237	3	1
9	1	78	0	0	0	0	0	58.57	24.200000	3	1

In [32]:

```
# this is to check the correlation
stroke_data.corr()
```

Out[32]:


	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	stroke
gender	1.000000	0.027404	-0.021390	-0.085749	-0.029366	-0.066212	-0.005529	-0.054213	0.025110	
age	0.027404	1.000000	0.276303	0.263668	-0.679181	-0.413852	-0.014169	0.238060	0.326271	
hypertension	-0.021390	0.276303	1.000000	0.108306	-0.164243	-0.075097	0.007913	0.174474	0.160189	
heart_disease	-0.085749	0.263668	0.108306	1.000000	-0.114644	-0.058263	-0.003092	0.161857	0.038899	
ever_married	-0.029366	-0.679181	-0.164243	-0.114644	1.000000	0.375944	0.006261	-0.155068	-0.335705	
work_type	-0.066212	-0.413852	-0.075097	-0.058263	0.375944	1.000000	-0.014573	-0.068135	-0.341973	
Residence_type	-0.005529	-0.014169	0.007913	-0.003092	0.006261	-0.014573	1.000000	0.004946	0.000120	
avg_glucose_level	-0.054213	0.238060	0.174474	0.161857	-0.155068	-0.068135	0.004946	1.000000	0.168751	
bmi	0.025110	0.326271	0.160189	0.038899	-0.335705	-0.341973	0.000120	0.168751	1.000000	
smoking_status	-0.039693	-0.376142	-0.129012	-0.064671	0.303694	0.307774	-0.004656	-0.097732	-0.230553	
stroke	-0.009200	0.245128	0.127904	0.134914	-0.108340	-0.058224	-0.015458	0.131945	0.038947	

In [33]:

```
plt.figure(figsize=(25,15))
sns.heatmap(stroke_data.corr(),annot=True)
```

Out[33]:

&lt;AxesSubplot:&gt;



## Splitting the dataset into dependant and independant variables

In [34]:

```
X = stroke_data.drop(['stroke'], axis=1)
y = stroke_data['stroke']
```

## Class Distribution

In [35]:

```
# Checking the class distribution of the target variable
class_dist = pd.Series(y).value_counts(normalize=True)
print("Class distribution:\n{}".format(class_dist))
```

Class distribution:

0 0.951272

1 0.048728

Name: stroke, dtype: float64

## Creating a baseline prediction

In [36]:

```
# Creating a baseline prediction that always predicts the majority class (0) for the test set
baseline_pred = np.zeros(len(y))

# Calculating the accuracy of the baseline prediction for the test set
baseline_acc = sum(baseline_pred == y) / len(y)

print("Baseline accuracy: {:.2f}%".format(baseline_acc * 100))
```

Baseline accuracy: 95.13%

## Machine Learning Models

### 1: Logistic Regression

In [37]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import cross_val_score
```

In [38]:

```
# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, random_state=0)
```

In [39]:

```
# Checking the class distribution of the target variable
class_dist = pd.Series(y_train).value_counts(normalize=True)
print("Class distribution:\n{}".format(class_dist))
```

```
Class distribution:
0    0.9512
1    0.0488
Name: stroke, dtype: float64
```

In [40]:

```
# Checking the class distribution of the target variable
class_dist = pd.Series(y_test).value_counts(normalize=True)
print("Class distribution:\n{}".format(class_dist))
```

```
Class distribution:
0    0.951487
1    0.048513
Name: stroke, dtype: float64
```

In [41]:

```
# Standardising the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Normalising the data
normalizer = MinMaxScaler()
X_train_norm = normalizer.fit_transform(X_train_scaled)
X_test_norm = normalizer.transform(X_test_scaled)
```

#### Dummy classifier to perform baseline accuracy score

In [42]:

```
from sklearn.dummy import DummyClassifier
dummy_clf = DummyClassifier(strategy="most_frequent")
dummy_clf.fit(X_train_scaled, y_train)
DummyClassifier(strategy='most_frequent')
#dummy_clf.predict(X_test_scaled)
dummy_clf.score(X_test_scaled, y_test)
```

Out[42]:

```
0.9514866979655712
```

## Training the Logistic Regression Classifier on the given training dataset

In [43]:

```
# Creating a logistic regression model object
logreg = LogisticRegression(penalty = "l2", C = 10, solver = "newton-cg")

# Fitting the model on the training data
logreg.fit(X_train_scaled, y_train)
```

Out[43]:

```
LogisticRegression(C=10, solver='newton-cg')
```

## Testing the Logistic Regression Classifier on testing dataset

In [44]:

```
# Predicting the labels of the test data
y_pred = logreg.predict(X_test_scaled)

# Printing the accuracy score of the model
print("Accuracy of the model on the test set is : {:.2f}".format(logreg.score(X_test_scaled, y_test)))
```

Accuracy of the model on the test set is : 0.95

In [45]:

```
#Performing the k-fold cross-validation on the LR model
# here we define k=5
scores = cross_val_score(logreg, X, y, cv=5)

# Printing the accuracy for each fold and the mean accuracy across all folds
for i, score in enumerate(scores):
    print("Fold {} accuracy: {:.2f}%".format(i+1, score * 100))
print("Mean accuracy: {:.2f}%".format(scores.mean() * 100))
print("Standard deviation: {:.2f}%".format(scores.std() * 100))
```

```
Fold 1 accuracy: 95.11%
Fold 2 accuracy: 95.11%
Fold 3 accuracy: 95.11%
Fold 4 accuracy: 95.01%
Fold 5 accuracy: 95.21%
Mean accuracy: 95.11%
Standard deviation: 0.06%
```

## The ROC (Receiver Operating Characteristic) curve graphical representation

In [46]:


```
from sklearn.metrics import roc_curve, roc_auc_score

# Predicting the probability scores of the test data
y_prob = logreg.predict_proba(X_test_norm)

# Computing the false positive rate, true positive rate, and threshold values
fpr, tpr, thresholds = roc_curve(y_test, y_prob[:,1])

# Computing the AUC score
auc_score = roc_auc_score(y_test, y_prob[:,1])

# Creating a ROC curve plot
plt.plot(fpr, tpr, color='darkorange', label=f'ROC curve (AUC = {auc_score:.2f})')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Logistic Regression (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



## 2. Support Vector Machine (SVM)

In [47]:

```
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
```

### Training the SVM Classifier on the given training dataset

In [48]:

```
# Training the SVM classifier on the standardised data
svm = SVC(kernel='linear', C=1.0, random_state=0)
svm.fit(X_train_scaled, y_train)
```

Out[48]:

```
SVC(kernel='linear', random_state=0)
```

### Testing the SVM Classifier on testing dataset

In [49]:

```
# Evaluating the classifier on the testing data
score = svm.score(X_test_scaled, y_test)
print("Accuracy of SVM Classifier on Testing Dataset:", score)
```

```
Accuracy of SVM Classifier on Testing Dataset: 0.9514866979655712
```

In [50]:

```
# Performing k-fold cross-validation on the SVM model
scores = cross_val_score(svm, X, y, cv=5)

# Printing the accuracy for each fold and the mean accuracy across all folds
for i, score in enumerate(scores):
    print("Fold {} accuracy: {:.2f}%".format(i+1, score * 100))
print("Mean accuracy: {:.2f}%".format(scores.mean() * 100))
print("Standard deviation: {:.2f}%".format(scores.std() * 100))
```

```
Fold 1 accuracy: 95.11%
Fold 2 accuracy: 95.11%
Fold 3 accuracy: 95.11%
Fold 4 accuracy: 95.11%
Fold 5 accuracy: 95.21%
Mean accuracy: 95.13%
Standard deviation: 0.04%
```

## Plotting the SVM ROC representation

In [51]:

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Computing the false positive rate (fpr) and true positive rate (tpr)
fpr, tpr, thresholds = roc_curve(y_test, svm.decision_function(X_test_scaled))

# Computing the area under the ROC curve (AUC)
roc_auc = auc(fpr, tpr)

# Plotting the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(' SVM (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

